

Database Searching & Indexing - Notes

Database Searching & Indexing

Basic Concepts

- **Record:** A collection of values for attributes of a single entity.
- **Collection:** A set of records representing the same entity type.
- **Search Key:** A single value for an attribute from an entity type.
- **Storage:**
 - Information is stored in records of X bytes across N memory locations.
- **Contiguous Allocation:**
 - Data stored in a continuous memory block.
- **Linked List:**
 - Each record stores X bytes + space for memory addresses.
 - Records are linked via memory addresses.
 - **Pros:** Fast insertion anywhere in the list.
 - **Cons:** Slower for random access.
- **Arrays:**
 - **Pros:** Faster for random access.
 - **Cons:** Slow insertions since other records need shifting.

Searching Algorithms

Binary Search

- Divides the dataset in half each step (left-biased).
- **Time Complexity:**
- **Best Case:** $O(1)$

- Worst Case: $O(\log n)$

Linear Search

- Sequentially checks each record.
- Time Complexity:
- Best Case: $O(1)$
- Worst Case: $O(n)$

Database Searching

- Data is stored on disk, indexed by ID values.
- Data cannot be sorted simultaneously by multiple attributes.
- Indexes help improve search performance.

Binary Search Trees (BST)

Example of a BST:

23

/ 17 31

/ \ 14 20 90

Tree Traversal Methods:

- Pre-order
- Post-order
- In-order
- Level-order traversal

AVL Trees (Self-Balancing BSTs)

- A BST that maintains balance using rotations.

- **Rotation Cases:**
- **LL (Left-Left) -> Single Right Rotation.**
- **RR (Right-Right) -> Single Left Rotation.**
- **LR (Left-Right) -> Double Rotation.**
- **RL (Right-Left) -> Double Rotation.**

Memory Hierarchy & Optimization

Key Concepts

- **Minimizing Disk Accesses:**
- **HDD/SSD storage is slow.**
- **We want to reduce secondary storage accesses.**
- **Key-Value Pointer:**
- **Uses 64-bit (8-byte) integers for pointers.**
- **SSD/HDD:**
- **Persistent (survives power loss).**
- **Reads in blocks, not individual bytes.**

B-Trees & B+ Trees

B-Trees

- **Balanced search tree for efficient indexing.**
- **Internal Nodes:**
- **Store keys and pointers to children.**
- **Leaf Nodes:**
- **Store actual data values.**

B+ Trees

- **Optimized for disk-based indexing.**

- **Properties:**
- **Root node does not need to be full.**
- **Insertions always happen at the leaf level.**
- **Minimizes disk accesses.**
- **Leaves are stored in a Doubly Linked List (DLL).**
- **Keys in nodes are kept sorted.**

Beyond Relational Databases (RDBMS)

Transactions

- **A sequence of CRUD operations executed as one unit.**
- **ACID Properties:**
 1. **Atomicity: All-or-nothing execution.**
 2. **Consistency: Transactions maintain a valid DB state.**
 3. **Isolation: No interference between concurrent transactions.**
 4. **Durability: Committed transactions persist.**

Concurrency Issues

- **Dirty Read: A transaction reads uncommitted data.**
- **Non-Repeatable Read: Data changes mid-transaction.**
- **Phantom Read: New rows appear or disappear mid-transaction.**

Scaling Strategies

- **Vertical Scaling: Adding more CPU/RAM to one machine.**
- **Horizontal Scaling: Adding more machines to handle the load.**
- **Distributed Systems:**
 - **A group of independent machines acting as one system.**
 - **No shared global clock.**

- **Sharding: Splitting data across nodes.**
- **Replication: Keeping multiple copies of data.**

CAP Theorem

- **In distributed databases, you can only have two of:**
 1. **Consistency (C): All nodes return the latest data.**
 2. **Availability (A): Every request gets a response.**
 3. **Partition Tolerance (P): The system functions despite network failures.**

Concurrency Control

Pessimistic Concurrency

- **Transactions lock data before reading/writing.**
- **Best for high-conflict environments (e.g., banking).**

Optimistic Concurrency

- **Transactions assume low conflicts, using timestamps.**
- **Best for low-conflict environments.**

NoSQL Databases

BASE Model (NoSQL Alternative to ACID)

- **Basically Available: System works most of the time.**
- **Soft State: The system can change without input.**
- **Eventual Consistency: Updates will propagate eventually.**

Types of NoSQL Databases

1. **Key-Value Stores (e.g., Redis)**
2. **Graph Databases (e.g., Neo4j)**

3. Columnar Databases (e.g., Apache Cassandra)

4. Document Stores (e.g., MongoDB)

5. Vector Databases (e.g., Pinecone)

Redis (Key-Value Store)

- **Keys: Strings.**
- **Values: Strings, lists, sets, hashes, sorted sets.**
- **Very fast and scalable.**
- **Stores frequently accessed data.**

Final Thoughts

- **RDBMS: Strong consistency, well-structured queries (SQL).**
- **NoSQL: Flexible, scalable, and optimized for distributed systems.**
- **B-Trees & Hash Tables: Key for efficient search & indexing.**
- **Distributed Systems: Balancing availability, consistency, and fault tolerance.**